



# STRIDE

*A personal reasoning agent over a self-logged life*

## THE PROBLEM

- We're bad at connecting our own dots — sleep, meals, movement, and mood interact over weeks, but we rarely see how.
- Existing trackers show numbers. They don't reason about them.

## THE THESIS — NOT JUST A HABIT TRACKER

- Two building blocks: Inputs (things you do or control) and Outcomes (things you observe).
- The LLM reasons over those statistical sketches + plain-language descriptions — surfacing grounded insights.

## HIGH-LEVEL DESIGN



- Deterministic math stays in SQL — means, trends, streaks, outliers. The LLM does not touch raw numbers (not a great choice for time series).
- Per-user inputs and outputs.
- Ease of logging is important: auto-extraction of data points from voice or chat descriptions.

## WHERE THIS GOES

### Per-user, eventually self-extending: from tracker to agent



(a) Each person's schema is their own; the near-term frill is letting a user define a new input or outcome widget in a few taps rather than waiting on us. On the input side specifically, Apple Health is the natural first integration — sleep, steps, and heart data flowing in without manual logging.

(b) Proposing new things to track, with the schema able to auto-evolve in the future, with minimal friction to the user.

### A more general framework than health



- Strip away the domain, and this is really an input-vs-outcome tracking-and-reasoning engine — statistical sketches in, grounded LLM reasoning out.
- The Wall, in principle, becomes more than bookmarks: a lightweight knowledge base — a stripped-down NotebookLM — that grounds the agent's reasoning in the person's own reference material, not just their logged sketches.
- Examples: a nutritionist's article pasted into the Wall informs how the agent reasons about meal timing and weight; a personal-finance version of Stride reasons over spending/savings inputs against balance/credit-score outcomes, grounded in the person's own budget notes.